

PHASE 1: DATA COLLECTION

The first phase focuses entirely on building the underlying (intelligent) investor infrastructure.

STEP 1. Taking in data from various sources

1. FundingStack

<https://fundingstack.com/>

Cost

\$250/user/month

Features

- Run unlimited private deals / companies / projects under a single account
- Analytics and activity reports on each deal
- Unlimited access to database
- **Access to 227,000+ global investors with contact information**
- **Get Intro feature**
- Ability to find warm introductions to investors
- 1500 investors per pipeline
- Multiple pipeline creation
- Unlimited investor updates with view tracking
- Pitch deck hosting with view tracking
- Email integration
- Personalized bulk email
- Open tracking
- **Virtual Data Room (\$150 add-on)**
- Docs & Templates collection
- Customer support
- **CRM export to CSV**

Essentially, we are going to use FundingStack exports to build our own sovereign allocator database via their CRM->CSV feature.

2. REGIONAL OPEN CSVs

In addition to FundingStack, we will ingest:

- regional angel investor lists

- syndicate databases
- venture firm directories
- European allocator lists
- Indian startup investor datasets

These datasets help reduce dependency on expensive proprietary systems while improving geographic coverage.

STEP 2. REGULATORY AUDIT GATE

Once raw records are ingested, they move through regulatory verification.

This layer is critical.

Most investor databases contain:

- stale firms
- dormant entities
- inactive allocators
- outdated contacts
- unverifiable structures

We want to verify active investors before outreach.

STEP 3. INTENT

Verification alone is not enough.

We also need to determine whether an allocator is actually aligned with our investment thesis.

This layer identifies:

- If they align to AI investment
- How often they back emerging manager/1st time fund raisers
- Silicon Valley adjacency
- YC ecosystem overlap

Emerging Manager Filter

We will flag allocators that participated in:

- syndicates
- rolling funds
- sub-\$50M fund closes

- early-stage deployments

AI Investment

The system checks whether an allocator has exposure to:

- OpenAI
- Anthropic
- Cohere
- xAI/Grok
- SpaceX
- mobility infrastructure
- Other

If an allocator has previously backed companies in these ecosystems, we may mark a label like

ai_investment_signal = TRUE

YC Exposure Mapping

We can cross-reference

- YC Demo Day participation
- YC founder networks
- mentor ecosystems
- Silicon Valley adjacency

This is important because:

- 60% of the broader network sits inside Silicon Valley
- warm graph proximity matters heavily in fundraising
- YC exposure creates immediate contextual trust (GIVE MONEY ON TRUST)

STEP 4. SMTP Handshake Validation

TO TRY TO MITIGATE BOUNCE RATES.

Suppose FundingStack lists:

Virat Kohli
Partner at XYZ Family Office

The system then:

1. discovers Kohli moved firms 6 months ago
2. validates the email server
3. performs SMTP handshake verification
4. identifies inbox failure or mismatch
5. flags the record as stale
6. removes the allocator from active outreach pools

Without this layer:

- bounce rates increase
- sender reputation deteriorates
- investor targeting quality collapses

Trying to set-up a back-end technical layer that verifies if an email address is correct, phone-number is correct, etc.

DATABASE**Supabase**

<https://supabase.com/>

Cost: Approximately \$25/month

How the database will look right now (END OF PHASE 1)

Suppose our data finds
ABC Family Office
Singapore
Backed Cohere indirectly
Participated in a YC Demo Day
Principal moved funds 2 years ago
Verified through MAS registry
Valid email confirmed

All this information will get stored in a structured investor pipeline. This pipeline will become usable by AI ranking systems, outreach mechanisms, relationship intelligence, personalization systems, etc.

We will essentially have a verified allocator database.

Example:

- 2,500 high-conviction allocator profiles
- verified investors
- family offices
- LPs
- syndicates
- emerging-manager allocators

MONTHLY COST STRUCTURE

Infrastructure Component	Estimated Cost
FundingStack	\$250/month
Supabase	~\$25/month
SEC EDGAR	Free
UK Companies House	Free
MAS Registry	Free
SEBI & IFSCA	Free
Regional Open CSVs	Free

Total Estimated Monthly Cost

Approximately \$275/month

PHASE 2: Investor Database

We are going to convert verified Phase 1 records into a searchable, structured, and AI-ready investor database. Every tag, filter, and classification built here becomes the foundation on which the AI agent, outreach engine, and pipeline system operate.

- investor organization
- tagging
- categorization
- filtering
- searchability
- relationship notes
- pipeline status
- operational usability

In phase 1 we found our allocators, and asked: is this allocator real and relevant?

In phase 2 we ask: how do we operationally manage and use this allocator?

1. Investor Organization: Every allocator record is classified into one of eight principal types:

- Family Offices
- Multi-Family Offices
- Angel Investors
- High Net Worth Individuals (HNIs)
- Fund of Funds
- Corporate Venture Capital
- Institutional Allocators
- Accelerators & Incubators

Although each type behaves differently, we're going to maintain a singular database. Our AI agent needs to query one place. If allocator types lived in separate tables, the agent would need to know where to look before it could do anything. However, a single unified system means one query returns all relevant allocators regardless of type, with type acting as a mere filter later on.

2. Database Architecture: We are using a two-layer architecture connected by a single allocator ID. The allocator data has two fundamentally different natures:

- a. **Static** : name, organization, geography, investor type, check size. This won't change most of the time, and when it does it'll a more significant event (fund closes, someone moves firms)
- b. **Dynamic** : ai_investment_signal, yc_exposure, emerging_manager_filter, email_verified, pipeline_status. These are going to change constantly, with every outreach, every new investment discovered, every update.

ALLOCATOR CORE TABLE

allocator_id	unique identifier
name	individual contact name
organization	firm or family office name
investor_type	One of the 7 principal types
geography	country and city
check_size_min	minimum investment size
check_size_max	maximum investment size
primary_contact	Email, phone, LinkedIn
verification_status	verified, unverified
data_added	when the record entered the system

ALLOCATOR SIGNALS TABLE

allocator_id	will link back to core table
signal_type	what kind of intelligence this is
signal_value	TRUE/FALSE/score/label
signal_date	when this signal was captured
source	where the signal came from
confidence_score	how reliable this signal is

For example; ABC Family Office, Singapore:

Core table will contain; ABC Family Office, Singapore, Family Office, Verified.

Signals table:

abc_001	ai_investment_signal	TRUE	Cohere portfolio match	0.91
---------	----------------------	------	------------------------	------

abc_001	yc_exposure	TRUE	Demo Day 2024 attendance	0.87
abc_001	emerging_manager_flag	TRUE	sub-\$40M fund close 2023	0.94

After this, Phase 3 will discover a new signal; For example, a Podcast in which AI has been discussed. This may become a new role in the signals table, however, the core record is untouched.

3. TAGGING & CATEGORIZATION

Identity Tags

- investor_type (Family Office, HNI, Fund of Funds, etc.)
- geography (country, city, region)
- check_size_range (micro / small / mid / large)
- fund_stage_preference (pre-seed / seed / Series A / multi-stage)

Thesis Signal Tags

- ai_investment_signal (TRUE / FALSE)
- yc_exposure (TRUE / FALSE)
- emerging_manager_appetite (TRUE / FALSE)
- silicon_valley_adjacency (TRUE / FALSE)

Verification Tags

- entity_verified (TRUE / FALSE)
- email_verified (TRUE / FALSE)
- regulatory_source (SEC / MAS / Companies House / SEBI)
- record_freshness (active / stale / unknown)

Relationship Tags

- warm_intro_available (TRUE / FALSE)
- intro_path (via FundingStack Get Intro / via network / unknown)
- prior_contact (TRUE / FALSE)
- last_interaction_date

4. FILTERING & SEARCHABILITY

It should be *queryable* in seconds. For example, we open on Monday and ask:

>> Show me all the verified Singapore family offices with ai_investment_signal=TRUE and an emerging manager appetite, sorted by check size.

Primary Filters:

- Investor type
- Geography (region / country / city)
- Check size range
- Thesis alignment (AI, mobility, early-stage)
- Verification status
- Warm intro availability
- Pipeline status

Compound Filters:

- investor_type = HNI AND geography = HK AND ai_investment_signal = TRUE

5. PIPELINE STATUS

Every allocator record carries a pipeline position and a relationship log.

1. Lead Sourced 2. Researched 3. Qualified 4. Warm Intro Identified 5. Outreach Initiated 6. First Meeting 7. Follow-Up 8. Due Diligence 9. Commitment 10. Closed LP

If meetings are transcribed, we can have freeform notes which can be interpreted.

PHASE 3

Deploy 6 specialized AI agents that operate continuously on top of Phase 2's database to turn static allocator data into living, actionable intelligence.

Phase 2's allocator_core table (2,500 allocator records)

Runtime: Continuous (daily + weekly cycles)

Total Monthly Cost: ~\$2.60 for all AI agents

Will be using OpenClaw (self-hosted, node-based DAG workflow)

THE 6 AGENTS (IN EXECUTION ORDER)

1. RESEARCH AGENT

Model: Gemini 2.5 Flash (primary), Claude Haiku 4.5 (fallback)

Trigger: Daily at 9 AM IST (cron)

Batch Size: 500 allocators per run (5 runs/day for 2,500 total)

Processing Time: ~12 minutes per batch

What It Does:

- Scans LinkedIn, podcasts, interviews, conference talks, news, portfolio announcements
- Extracts signals, ignores noise
- Writes findings directly into signals table

Output Written to Supabase:

- Recent public activity log
- New investment signals
- Network movement flags (firm change, fund close, new role)
- Updated thesis intelligence
- Confidence scores on all signals

Optimization: Batch API (50% cost reduction) + Prompt Caching (90% on repeat profiles)

Monthly Cost: ~\$0.30

2. INTENT AGENT

Model: Gemini 2.5 Flash (primary), Claude Haiku 4.5 (fallback)

Trigger: After Research Agent completes (node completion)

Batch Size: 2,500 allocators (single batch via prompt caching)

Processing Time: ~15 minutes for full batch

What It Does:

- Finds who is ready to move right now (not just thesis-aligned)
- Looks for recent fund closes, new LP commitments, Demo Day attendance, syndicate participation
- Splits every allocator into two groups:
- Group A: Aligned and deployable → contact now
- Group B: Aligned but dormant → monitor, contact later

Output Written to Supabase:

intent_score (0–100)

deployment_signal (active, dormant, unknown)

last_deployment_event with date

recommended_timing (contact now, monitor, deprioritise)

Optimization: Prompt Caching (90% cost reduction after first run)

Monthly Cost: ~\$0.20

3. RANKING AGENT

Model: Claude Haiku 4.5 (primary), Gemini 2.5 Flash (fallback)

Trigger: Monday at 9 AM IST (cron, weekly)

Batch Size: 2,500 allocators (single batch)

Processing Time: ~20 minutes for full ranking

What It Does:

Every allocator gets a conviction_score that updates continuously

It's not static. allocator who attended YC Demo Day last week moves up, dormant fund close moves them down

Will try to always reflect current market state

Scoring Formula:

$$\text{conviction_score} = 0.4 \times \text{fit_score} + 0.3 \times \text{accessibility_score} + 0.3 \times \text{urgency_score}$$

Output Written to Supabase:

conviction_score (0–100) per allocator (stored in allocator_core)

weekly_priority_list of top 20 (stored in weekly_rankings table)

score_delta since last run

one-sentence reason per allocator

Optimization: Batch API (50% cost reduction)

Monthly Cost: ~\$0.80

4. PERSONALIZATION AGENT

Model: Claude Sonnet 4.6 (primary), Claude Haiku 4.5 (fallback)

Trigger: After Ranking Agent completes (node completion)

Batch Size: Top 20 allocators (1 per week)

Processing Time: ~5 minutes for 20 profiles

What It Does:

- MOST IMPORTANT AGENTS; produces what LP actually reads
- Generates warm_thread (2–3 sentence contextual hook, not template)
- Generates full outreach email in fund's voice, ready to review
- Generates 2-sentence LinkedIn DM version
- Recommends channel (email, LinkedIn DM, warm intro)

Output Written to Supabase:

warm_thread (2–3 sentence contextual hook)

full_draft_email in fund's voice

draft_dm (2-sentence LinkedIn version)

Recommended_channel

Optimization: Prompt Caching on fund context (90% cost reduction)

Monthly Cost: ~\$0.60

5. CONNECT AGENT

Model: Claude Haiku 4.5 (primary), Gemini 2.5 Flash (fallback)

Trigger: Monday at 9 AM IST (cron, weekly, after Personalization Agent)

Batch Size: 70 LinkedIn requests/week

Processing Time: ~10 minutes for queue creation

What It Does:

- Manages LinkedIn sequencing
- Reads weekly_priority_list from Ranking Agent
- Pulls DM drafts from Personalization Agent
- Sequences for review
- Tracks every request sent, accepted, ignored, declined
- Schedules follow-ups
- Flags accepted connections with no reply as warm signals

LinkedIn Limits:

- 20–25 requests per day
- 70 per week hard cap (below LinkedIn's 100 ceiling)
- 14-day manual warm-up before automation begins
- Timing randomized within business hours only (9 AM–6 PM, Mon–Fri)

Output Written to Supabase:

weekly_connect_queue of 70 approved requests sequenced

Follow-up triggers for accepted connections with no reply

Cadence health report (account safety status)

Response tracking per allocator (accepted, pending, declined)

Optimization: Batch API (50% cost reduction)

Monthly Cost: ~\$0.30

WORKFLOW SEQUENCE (DAG) [Directed Acyclic Graph]

Weekly Pipeline (Monday 9 AM IST):

1. Research Agent (Daily at 9 AM)

↓ Scans 2,500 allocators, writes signals to Supabase

2. Intent Agent (After Research Agent completes)

↓ Evaluates deployment signals, writes intent_score to Supabase

3. Ranking Agent (Monday 9 AM)

↓ Scores all 2,500 allocators, writes conviction_score + top 20 to Supabase

4. Personalization Agent (After Ranking Agent completes)

↓ Generates warm_thread + drafts for top 20, writes to Supabase

5. Connect Agent (Monday 9 AM, after Personalization Agent)

↓ Creates weekly_connect_queue (70 requests), writes to Supabase

6. Outreach Automation (Phase 4, triggered by Connect Agent)

↓ DM Sequencer sends LinkedIn requests + email sequences

7. Loop back to Research Agent (daily cycle continues)

TECHNICAL DEPENDENCIES

SUPABASE (https://supabase.com/)	DATABASE
OpenClaw (https://openclaw.ai/)	ORCHESTRATION LAYER (DAG WORKFLOW CONNECTING ALL AGENTS)
Apollo.io API	Phase 2 Contact Data (Email, Phone, LinkedIn URL)
LinkedIn Sales Navigator	LinkedIn data +connection requests (Connect Agent)
Claude API	Haiku 4.5 + Sonnet 4.6 Models
Gemini API	2.5 Flash model

DELIVERABLES AT END OF PHASE 3

1. Live DAG workflow in OpenClaw connecting all 6 agents
2. Populated signals table with dynamic research + intent signals
3. Weekly conviction_scores for all 2,500 allocators
4. Automated weekly_priority_list of top 20 every Monday
5. Warm threads + draft emails ready for review (top 20)
6. Weekly LinkedIn queue of 70 approved requests

Total AI cost: ~\$2.60/month